

ELOQUENT sensemaking

Conference and Labs of the Evaluation Forum - Information Access Evaluation
meets Multilinguality, Multimodality, and Visualization - CLEF 2025



OutstandingOutsiders

Given:

Task

→ **Context: English text** (large, up to 35k tokens)

1. Teacher (other teams):

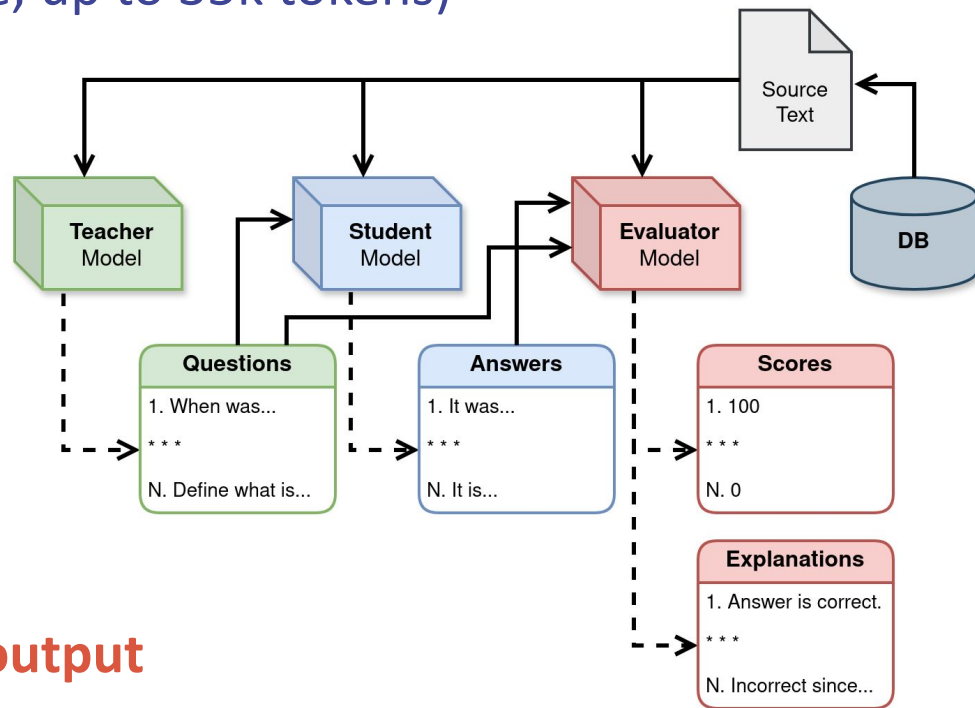
→ **Question output**

2. Student (other teams):

→ **Answer output**

3. Evaluator:

→ **Score (+ Explanation) output**



Why bother?

*“LLMs store immense amounts of information and easily answer questions about it. **Can they limit their knowledge to information provided in given materials?**”* - task provider.

- Potential help for real teachers and evaluators
- Testing of models knowledge constraints
- Exploring different aspects of source-based QA task

Problems to deal

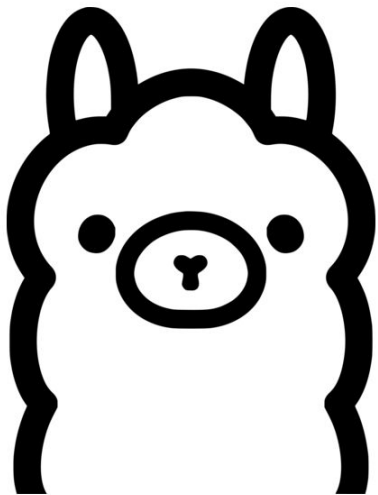
Too large context

Hard to fine-tune a specific model locally because:

- No available data with a comparable context-question size
- Models with such large context window take a lot of time to finetune
- Manual evaluation of predictions is exhausting



Solution



Use ollama models

- Various open source models to run inference (e.g. [deepseek](#), [llama](#) and [llava](#), [qwen](#), [gemma](#))
- Different parameter sizes available for some models
- Framework **optimized for multiple GPUs** using free video RAM as a buffer for input tokenization besides storing the model weights

Choosing models

Tech limitations:

- **GPU memory available** (eventually used **2 x NVIDIA A30** (24Gb) or **3 x NVIDIA RTX A4000** (16Gb) from UFAL cluster)
- **Context window size** (should be extended from the **ollama**'s cropped default)

General limitations:

- **JSON output format** (optional - can be post corrected)

Preliminary experiments

Goal: generate a training data for the Evaluator e.g. **Q,A,C-to-S** which is, in fact, **text-to-int** task

Data: QA dataset with answers including source quotes was parsed to **Question, Answer, Context** variables - then expanded with **wrong** (**prev/next** string in sequence) **Answers** in 1:1 ration to correct ones

Run: inference of **llama3.3** as Evaluator where input's **Context** is composed out of **prev**, **current** (actually relevant), and **next** items in the data sequence, and output prompted to be a JSON containing a **Score**

Preliminary experiments

Results:

- Some predictions with incorrect JSON formatting (e.g. missed punctuation) - can be fixed by regex
- Considering 1:1 ratio of correct-wrong **Answers**, the **Scores** were not so strict in proportions - **many initially right Answers scored zero**

Conclusions:

Decided to forget about creating a training dataset or using old llamas

Applied approach

Models used (**128k context window size**):

Qwen3 - 30b - a3b

MoE type transformer

Thinking/Not-thinking modes

Gemma3 - 27b

Decoder-only type transformer

Sliding Window Attention

Function calling head (structured output)



Qwen



Gemma

Prompt engineering



Decisions to make (based on a devset):

- **Define order of the given information** (e.g. task or Inputs first)
- **Describe task definition** (“You are an evaluator...”)
- **Give output format limitations** (e.g. JSON with exactly 2 fields)
- **State knowledge boundaries** (“Given the current context...”)

System prompt:

You are a fair teacher who grades students' answers. Evaluate the quality of the **Answer** specifically in response to the **Question** considering the **Context** provided.

Format your entire response as a single JSON object containing 'score' (an integer between 0 and 100, where 100 is best) and 'explanation' (a string briefly justifying the score).

Changeable prompt:

Question: {question}

Answer: {answer}

And given the following context: {text_fragments}

Generated output:

Potential thinking...

```
{"score": {int}, "explanation": {str}}
```

Applied approach

System prompt:

You are a fair teacher who grades students' answers. Evaluate the quality of the **Answer** specifically in ... a 'score' (an integer between 0 and 100, where 100 is best) and 'explanation' (a string briefly justifying the score).

Changeable prompt:

Question: What is a result of misplacing punctuation marks in machine translation?

Answer: What is a result of misplacing punctuation marks in machine translation is not discussed in the provided text.

And given the following context: Computation Graphs For our example neural network from Section ... book published in year 2020.

Generated output:

The text talks about ... thus the score is 100.

```
{"score": 100, "explanation": "The given answer is correct, the text does not mention misplacing punctuation marks in machine translation."}
```

Example

System prompt:

You are a fair teacher who grades students' answers. Evaluate the quality of the **Answer** specifically in ... a 'score' (an integer between 0 and 100, where 100 is best) and 'explanation' (a string briefly justifying the score).

Changeable prompt:

Question: What was the structure of trade in the Roman Empire?

Answer: In hospitals, dehydration is commonly treated with infusions.

And given the following context: CHAPTER OUTLINE 7.1 The Daily Life of a Roman Family 7.2 Slavery in the Roman Empire 7.3 The Roman Economy: Trade, Taxes, and Conquest ... Jewish population during the imperial period.

Generated output:

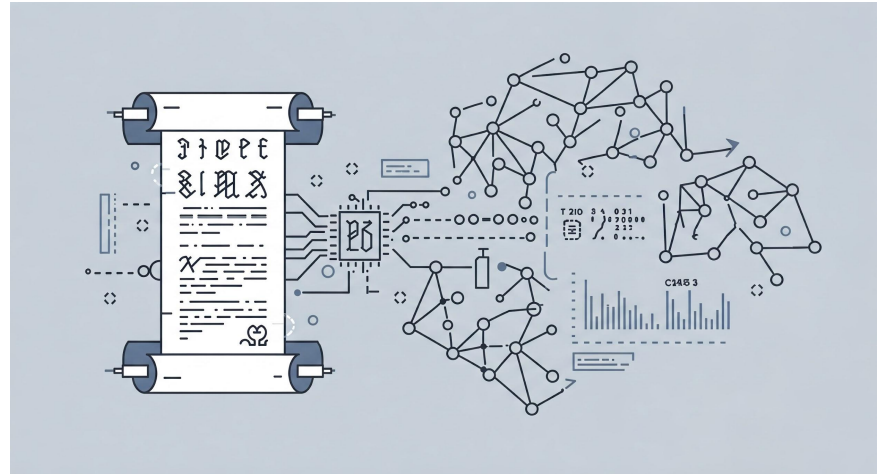
```
{"score": 10, "explanation": "The answer is entirely unrelated to the question. It is the full text of a chapter on the Roman Empire. There is no attempt to answer the question at all. Therefore the score is the lowest possible."}
```

Example

Findings

It's difficult to pick “the right” model when multiple of the same quality are available - too large and detailed context for a simple eyeball check of the generated outputs

Fine-tuning models is inefficient in comparison to prompt-engineering of the newest models (lack of time and equal in style training datasets)





Thanks for attention